


ENM531: Data-driven modeling and probabilistic scientific computing

Lecture #10: Physics-informed neural networks

Paris Perdikaris
February 20, 2020



SC15-009: Recent Advances in Physics-Informed Deep Learning



Instructors:

- Paris Perdikaris (UPenn, pgp@seas.upenn.edu)
- Maziar Raissi (NVIDIA, maziar.raissi@gmail.com)

Schedule (Room 205)

Time	Lecturer	Topic
8.30-9.20am	Paris Perdikaris	Supervised learning with neural networks in Tensorflow
9.20-10.10am	Maziar Raissi	Physics-informed neural networks (Part I)
10.10-10.30am	Coffee Break	
10.30-11.20am	Paris Perdikaris	Physics-informed neural networks (Part II)
11.20-12.10pm	Maziar Raissi	Multi-step neural networks
12.10-1.00pm	Lunch Break	
1.00-1.50pm	Paris Perdikaris	PINNs on Graphs
1.50-2.20pm	Maziar Raissi	Hidden physics models
2.20-3.10pm	Paris Perdikaris	Physics-informed deep generative models
3.10-3.30pm	Coffee Break	
3.30-4.20pm	Maziar Raissi	Forward Backward Stochastic Neural Networks
4.20-5.10pm	Paris Perdikaris	Open challenges
5.10-5.30pm	Maziar Raissi	Summary and future work

Motivation and open challenges

Goal: Predictive modeling, analysis and optimization of complex systems



Challenges:

- High cost of data acquisition
- Limited and high-dimensional data
- Multiple tasks and data modalities (e.g. images, time-series, scattered measurements, etc.)
- Large parameter spaces
- Incomplete models, imperfect data (e.g., missing data, outliers, complex noise processes)
- Uncertainty quantification
- Robust design/control



Hypothesis:

- Can we bridge knowledge from scientific computing and machine learning to tackle these challenges?



Physics of AI: Two schools of thought

1. Physics is implicitly baked in specialized neural architectures with strong inductive biases (e.g. invariance to simple group symmetries).



*figures from Kondor, R., Son, H.T., Pan, H., Anderson, B., & Trivedi, S. (2018). Covariant compositional networks for learning graphs. arXiv preprint arXiv:1801.02144.

2. Physics is explicitly imposed by constraining the output of conventional neural architectures with weak inductive biases.


Psichogios & Ungar, 1992

Lagaris et. al., 1998

Raissi et. al., 2019


Lu et. al., 2019

Zhu et. al., 2019



Physics-informed Neural Networks

$$f \left(\mathbf{x}; \frac{\partial u}{\partial x_1}, \dots, \frac{\partial u}{\partial x_d}; \frac{\partial^2 u}{\partial x_1 \partial x_1}, \dots, \frac{\partial^2 u}{\partial x_1 \partial x_d}; \dots; \boldsymbol{\lambda} \right) = 0, \quad \mathbf{x} \in \Omega, \quad \mathcal{B}(u, \mathbf{x}) = 0 \quad \text{on} \quad \partial\Omega,$$



Psichogios, D. C., & Ungar, L. H. (1992). A hybrid neural network–first principles approach to process modeling. *AIChE Journal*, 38(10), 1499-1511.

Lagaris, I. E., Likas, A., & Fotiadis, D. I. (1998). Artificial neural networks for solving ordinary and partial differential equations. *IEEE transactions on neural networks*, 9(5), 987-1000.

Raissi, M., Perdikaris, P., & Karniadakis, G. E. (2019). Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378, 686-707.

Lu, L., Meng, X., Mao, Z., & Karniadakis, G. E. (2019). DeepXDE: A deep learning library for solving differential equations. *arXiv preprint arXiv: 1907.04502*.

Physics-informed Neural Networks

Example: Burgers' equation in 1D

$$u_t + uu_x - (0.01/\pi)u_{xx} = 0, \quad x \in [-1, 1], \quad t \in [0, 1], \quad (3)$$
$$u(0, x) = -\sin(\pi x),$$
$$u(t, -1) = u(t, 1) = 0.$$


Let us define $f(t, x)$ to be given by

$$f := u_t + uu_x - (0.01/\pi)u_{xx},$$

```
def u(t, x):
    u = neural_net(tf.concat([t,x],1), weights, biases)
    return u
```

Correspondingly, the *physics informed neural network* $f(t, x)$ takes the form

```
def f(t, x):
    u = u(t, x)
    u_t = tf.gradients(u, t)[0]
    u_x = tf.gradients(u, x)[0]
    u_xx = tf.gradients(u_x, x)[0]
    f = u_t + u*u_x - (0.01/tf.pi)*u_xx
    return f
```



Physics-informed Neural Networks

The shared parameters between the neural networks $u(t, x)$ and $f(t, x)$ can be learned by minimizing the mean squared error loss

$$MSE = MSE_u + MSE_f, \quad (4)$$

where

$$MSE_u = \frac{1}{N_u} \sum_{i=1}^{N_u} |u(t_u^i, x_u^i) - u^i|^2,$$

and

$$MSE_f = \frac{1}{N_f} \sum_{i=1}^{N_f} |f(t_f^i, x_f^i)|^2.$$

Here, $\{t_u^i, x_u^i, u^i\}_{i=1}^{N_u}$ denote the initial and boundary training data on $u(t, x)$ and $\{t_f^i, x_f^i\}_{i=1}^{N_f}$ specify the collocations points for $f(t, x)$. The loss MSE_u corresponds to the initial and boundary data while MSE_f enforces the structure imposed by equation (3) at a finite set of collocation points.

Physics-informed Neural Networks






Figure 1: *Burgers' equation:* *Top:* Predicted solution $u(t, x)$ along with the initial and boundary training data. In addition we are using 10,000 collocation points generated using a Latin Hypercube Sampling strategy. *Bottom:* Comparison of the predicted and exact solutions corresponding to the three temporal snapshots depicted by the white vertical lines in the top panel. The relative \mathcal{L}_2 error for this case is $6.7 \cdot 10^{-4}$. Model training took approximately 60 seconds on a single NVIDIA Titan X GPU card.

Physics-informed Neural Networks




Physics-informed neural networks



Raissi, M., Yazdani, A., & Karniadakis, G. E. (2020). Hidden fluid mechanics: Learning velocity and pressure fields from flow visualizations. *Science*.


Recent advances

Discovery of ODEs




Raissi, M., Perdikaris, P., & Karniadakis, G. E. (2018). Multistep Neural Networks for Data-driven Discovery of Nonlinear Dynamical Systems. *arXiv preprint*

Discovery of PDEs




Raissi, M. (2018). Deep Hidden Physics Models: Deep Learning of Nonlinear Partial Differential Equations. *arXiv preprint arXiv:1801.06637*.

High-dimensional PDEs



Raissi, M. (2018). Forward-backward stochastic neural networks: Deep learning of high-dimensional partial differential equations. *arXiv preprint arXiv:1804.07010*.


Stochastic PDEs



Yang, Y., & Perdikaris, P. (2019). Adversarial uncertainty quantification in physics-informed neural networks. *Journal of Computational Physics*.


Recent advances

Fractional PDEs




Pang, G., Lu, L., & Karniadakis, G. E. (2018). fpinns: Fractional physics-informed neural networks. *arXiv preprint arXiv: 1811.08967*.

Surrogate modeling & high-dimensional UQ




Zhu, Y., Zabaras, N., Koutsourelakis, P. S., & Perdikaris, P. (2019). Physics-constrained deep learning for high-dimensional surrogate modeling and uncertainty quantification without labeled data. *Journal of Computational Physics*, 394, 56-81.

Multi-fidelity modeling for stochastic systems



Conditional deep surrogate models for stochastic, high-dimensional, and multi-fidelity systems. *Computational Mechanics*, 1-18.

Integrated software



Lu, L., Meng, X., Mao, Z., & Karniadakis, G. E. (2019). DeepXDE: A deep learning library for solving differential equations. *arXiv preprint arXiv:1907.04502*.